

Q3

① For the implementation of the k-medoids framework I chose the initial medoids randomly. Then I calculated the distance from every data point to the given medoids. Using the euclidean distance $d(p, q) = \sqrt{(p - q)^2}$ but also manhattan distance $d(p, q) = \sum_{i=1}^n |p_i - q_i|$. Afterwards the data points were assigned to the cluster with the least distance.

Following the medoids were adjusted, ~~that is~~ ~~it is~~ The center for every cluster was updated with the data point which was closest to the centroid of ~~the~~ each cluster. These steps were repeated for a certain amount of iterations or when the distortion function didn't decrease. Although, I mainly focused on the max iteration since the distortion function was causing performance issues.

② Depending on the number of the clusters k , a different amount of colours were displayed in the compressed pictures. ~~This also resulted in~~ The more clusters the more iterations until ~~convergence~~ convergence. The observation for different k where the following:

| #Cluster | 2 | 3 | 4 | 7 |
|------------|----|----|----|-------|
| #Iteration | 10 | 13 | 19 | >30 * |

* runtime error
manually stopped

③ Differences in the pictures based on different initial centroids were not significant. However, the run time for the clustering varied based on the different initial starts. When choosing a poor assignment as initial centroids (e.g. really close together) I was observing a considerable increase in terms of iterations and run time.

④ In general the solutions given through the k -means algo don't differ significantly from k -medoids. However, for smaller cluster values it appears that k -medoid is more effective capturing the separation of the colours. The running time was quite similar to k -medoids with a slight increase for the football picture (since it's bigger). This was quite surprising since k -medoids is supposed to be much faster.